

Activation Functions

- [ELU](#)
- [ReLU](#)
- [LeakyReLU](#)
- [Sigmoid](#)
- [Tanh](#)
- [Softmax](#)


ELU


Be the first to [contribute!](#)

ReLU

A recent invention which stands for Rectified Linear Units. The formula is deceptively simple: $\max(0, z)$. Despite its name and appearance, it's not linear and provides the same benefits as Sigmoid but with better performance.

Function	Derivative
$R(z) = \begin{cases} z & z > 0 \\ 0 & z \leq 0 \end{cases}$	$R'(z) = \begin{cases} 1 & z > 0 \\ 0 & z \leq 0 \end{cases}$





Pros

- Pro 1

Cons

- Con 1


Further reading


- [Deep Sparse Rectifier Neural Networks](#) Glorot et al., (2011)
- [Yes You Should Understand Backprop](#), Karpathy (2016)

LeakyReLU

LeakyRelu is a variant of ReLU. Instead of being 0 when $z < 0$, a leaky ReLU allows a small, non-zero, constant gradient α (Normally, $\alpha = 0.01$). However, the consistency of the benefit across tasks is presently unclear. ^[1]

Function	Derivative
$R(z) = \begin{cases} z & z > 0 \\ \alpha z & z \leq 0 \end{cases}$	$R'(z) = \begin{cases} 1 & z > 0 \\ \alpha & z \leq 0 \end{cases}$





Pros

- Pro 1

Cons

- Con 1


Further reading


- [Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification](#), Kaiming He et al. (2015)

Sigmoid


Sigmoid takes a real value as input and outputs another value between 0 and 1. It's easy to work with and has all the nice properties of activation functions: it's non-linear, continuously differentiable, monotonic, and has a fixed output range.

Function	Derivative
$S(z) = \frac{1}{1 + e^{-z}}$	$S'(z) = S(z) \cdot (1 - S(z))$





```
def sigmoid(z):
    return 1.0 / (1 + np.exp(-z))
```



```
def sigmoid_prime(z):
    return sigmoid(z) * (1-sigmoid(z))
```

Pros

- Pro 1

Cons

- Con 1


Further reading


- [Yes You Should Understand Backprop](#), Karpathy (2016)

Tanh


Tanh squashes a real-valued number to the range $[-1, 1]$. It's non-linear. But unlike Sigmoid, its output is zero-centered. Therefore, in practice the tanh non-linearity is always preferred to the sigmoid nonlinearity. [\[1\]](#)

Function	Derivative
$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	$\tanh'(z) = 1 - \tanh(z)^2$





```
def tanh(z):  
    return (np.exp(z) - np.exp(-z))  
           / (np.exp(z) + np.exp(-z))
```



```
def tanh_prime(z):  
    return 1 - np.power(tanh(z), 2)
```

Pros

- Pro 1

Cons

- Con 1

Softmax 🔗

Be the first to [contribute!](#)

References

- [1] (1, 2) <http://cs231n.github.io/neural-networks-1/>